

# Econ 425 Week 9

## Clustering and PCA

Grigory Franguridi

UCLA Econ

USC CESR

franguri@usc.edu

# Motivation

- what if a dataset has too many variables?
- e.g., most of the variables are correlated on analysis
- may lead to poor accuracy in estimation
- **dimension reduction** methods
  - Principal Component Analysis (PCA)

# Motivation

- what if a dataset has too many variables?
- e.g., most of the variables are correlated on analysis
- may lead to poor accuracy in estimation
- **dimension reduction** methods
  - Principal Component Analysis (PCA)

# Motivation

- what if a dataset has too many variables?
- e.g., most of the variables are correlated on analysis
- may lead to poor accuracy in estimation
- **dimension reduction** methods
  - Principal Component Analysis (PCA)

# Motivation

- what if a dataset has too many variables?
- e.g., most of the variables are correlated on analysis
- may lead to poor accuracy in estimation
- **dimension reduction** methods
  - Principal Component Analysis (PCA)

# Motivation

- what if a dataset has too many variables?
- e.g., most of the variables are correlated on analysis
- may lead to poor accuracy in estimation
- **dimension reduction** methods
  - Principal Component Analysis (PCA)

# PCA

- summarizes the information content in large datasets with a smaller dataset of “summary indices” that can be more easily visualized and analyzed
- underlying data can be measurements describing properties of production samples, chemical compounds or reactions, time points of a continuous process, batches from a batch process, biological individuals, or trials of a DOE protocol
- often used in the preliminary data analysis, before running any ML tasks

# PCA

- summarizes the information content in large datasets with a smaller dataset of “summary indices” that can be more easily visualized and analyzed
- underlying data can be measurements describing properties of production samples, chemical compounds or reactions, time points of a continuous process, batches from a batch process, biological individuals, or trials of a DOE protocol
- often used in the preliminary data analysis, before running any ML tasks



# PCA

- summarizes the information content in large datasets with a smaller dataset of “summary indices” that can be more easily visualized and analyzed
- underlying data can be measurements describing properties of production samples, chemical compounds or reactions, time points of a continuous process, batches from a batch process, biological individuals, or trials of a DOE protocol
- often used in the preliminary data analysis, before running any ML tasks

# How PCA works

- $X$  is data matrix with  $N$  rows (observations) and  $K$  columns (features)
- construct a variable space with as many dimensions as there are variables (see figure on the next slide)
- each variable represents a coordinate axis; for each variable, the length is standardized, typically by scaling to unit variance


# How PCA works

- $X$  is data matrix with  $N$  rows (observations) and  $K$  columns (features)
- construct a variable space with as many dimensions as there are variables (see figure on the next slide)
- each variable represents a coordinate axis; for each variable, the length is standardized, typically by scaling to unit variance

# How PCA works

- $X$  is data matrix with  $N$  rows (observations) and  $K$  columns (features)
- construct a variable space with as many dimensions as there are variables (see figure on the next slide)
- each variable represents a coordinate axis; for each variable, the length is standardized, typically by scaling to unit variance


## How PCA works



Feature space  $\mathbb{R}^K$ . Only three variable axes displayed. The “length” of each coordinate is standardized


# How PCA works

- each observation in  $X$  is a point in the feature space  $\mathbb{R}^K$




# How PCA works

- **centering**: subtract variable averages from the data. The vector of averages is the red point in  $\mathbb{R}^K$



# How PCA works

- subtraction of the average corresponds to a re-positioning of the origin of the coordinate system to the average point





# How PCA works: first principal component

- ready to compute the first principal component (PC1)
- PC1 is the line through the average point that **best approximates the data in the least squares sense**
- each observation (yellow dot) may now be projected onto this line to get the coordinate value along the PC-line (**score**)


# How PCA works: first principal component

- ready to compute the first principal component (PC1)
- PC1 is the line through the average point that **best approximates the data in the least squares sense**
- each observation (yellow dot) may now be projected onto this line to get the coordinate value along the PC-line (**score**)

# How PCA works: first principal component

- ready to compute the first principal component (PC1)
- PC1 is the line through the average point that **best approximates the data in the least squares sense**
- each observation (yellow dot) may now be projected onto this line to get the coordinate value along the PC-line (**score**)

## How PCA works: first principal component



- the first principal component (PC1) represents the *maximum variance direction* in the data


## How PCA works: second principal component

- usually one summary index or principal component is insufficient to model the systematic variation of data
- second principal component (PC2) is represented by a line through the average point *orthogonal* to the first PC

## How PCA works: second principal component


- usually one summary index or principal component is insufficient to model the systematic variation of data
- second principal component (PC2) is represented by a line through the average point *orthogonal* to the first PC

## The second principal component



- second principal component (PC2) reflects the second largest source of variation in the data while being *orthogonal* to the first PC
- PC2 also passes through the average point

## The second principal component



- second principal component (PC2) reflects the second largest source of variation in the data while being *orthogonal* to the first PC
- PC2 also passes through the average point



# How to calculate PC1 and PC2?

- **standardize the data**: each variable should be mean 0 and SD 1 (PCA is sensitive to scaling)
- calculate the sample covariance matrix

$$\hat{\Sigma} = \frac{1}{n-1} X^T X$$

- elements are covariances between each pair of features

# How to calculate PC1 and PC2?

- **standardize the data:** each variable should be mean 0 and SD 1 (PCA is sensitive to scaling)
- **calculate the sample covariance matrix**

$$\hat{\Sigma} = \frac{1}{n-1} X^T X$$

- elements are covariances between each pair of features

# How to calculate PC1 and PC2?

- **standardize the data:** each variable should be mean 0 and SD 1 (PCA is sensitive to scaling)
- **calculate the sample covariance matrix**

$$\hat{\Sigma} = \frac{1}{n-1} X^T X$$

- elements are covariances between each pair of features

# How to calculate PC1 and PC2?

- **calculate the eigenvalues and eigenvectors of the covariance matrix:** eigenvectors/eigenvalues represent directions of maximum variance / magnitude of variance in the data
- eigenvectors  $\{v_1, v_2, \dots, v_K\}$  and eigenvalues  $\{\lambda_1, \lambda_2, \dots, \lambda_K\}$  satisfy

$$\hat{\Sigma} v_k = \lambda_k v_k$$

- eigendecomposition:

$$\hat{\Sigma} = V \Lambda V^T,$$

where  $V$  is a matrix of eigenvectors,  $\Lambda$  is a diagonal matrix with eigenvalues on the diagonal

# How to calculate PC1 and PC2?

- **calculate the eigenvalues and eigenvectors of the covariance matrix:** eigenvectors/eigenvalues represent directions of maximum variance / magnitude of variance in the data
- eigenvectors  $\{v_1, v_2, \dots, v_K\}$  and eigenvalues  $\{\lambda_1, \lambda_2, \dots, \lambda_K\}$  satisfy

$$\hat{\Sigma} v_k = \lambda_k v_k$$

- eigendecomposition:

$$\hat{\Sigma} = V \Lambda V^T,$$

where  $V$  is a matrix of eigenvectors,  $\Lambda$  is a diagonal matrix with eigenvalues on the diagonal

## How to calculate PC1 and PC2?

- **calculate the eigenvalues and eigenvectors of the covariance matrix:** eigenvectors/eigenvalues represent directions of maximum variance / magnitude of variance in the data
- eigenvectors  $\{v_1, v_2, \dots, v_K\}$  and eigenvalues  $\{\lambda_1, \lambda_2, \dots, \lambda_K\}$  satisfy

$$\hat{\Sigma} v_k = \lambda_k v_k$$

- eigendecomposition:

$$\hat{\Sigma} = V \Lambda V^T,$$

where  $V$  is a matrix of eigenvectors,  $\Lambda$  is a diagonal matrix with eigenvalues on the diagonal

## How to calculate the PC1 and PC2 ?

- **sort the eigenvectors by their corresponding eigenvalues in descending order:** the eigenvector associated with the largest eigenvalue is the first principal component, and the eigenvector associated with the second largest eigenvalue is the second principal component:

$$PC1 = v_1$$

$$PC2 = v_2$$

# The model plane

- PC1 and PC2 together define a plane in  $\mathbb{R}^K$
- visualize the data by projecting observations onto this low-dimensional subspace and plotting (**score plot**)
- coordinate values of the observations on this plane **scores**




# The model plane

- PC1 and PC2 together define a plane in  $\mathbb{R}^K$
- visualize the data by projecting observations onto this low-dimensional subspace and plotting (**score plot**)
- coordinate values of the observations on this plane **scores**

## The model plane

- PC1 and PC2 together define a plane in  $\mathbb{R}^K$
- visualize the data by projecting observations onto this low-dimensional subspace and plotting (**score plot**)
- coordinate values of the observations on this plane **scores**

## The model plane



- PC1 and PC2 form a plane, which can be visualized graphically. Projections of observations onto the plane are called **scores**

## What is the score?

- PC scores of (standardized)  $X$  are obtained by multiplying  $X$  by the loadings (eigenvectors) of the covariance of  $X$ , say  $\hat{\Sigma}$
- recall that  $V$  is the matrix of eigenvectors (loadings) of  $\hat{\Sigma}$
- order the columns of  $V$  by their corresponding eigenvalues in descending order
- scores:

$$T_{n \times k} = XV$$

- the first column of  $T$  contains the scores for PC1, the second column contains the scores for PC2, etc.

# What is the score?

- PC scores of (standardized)  $X$  are obtained by multiplying  $X$  by the loadings (eigenvectors) of the covariance of  $X$ , say  $\hat{\Sigma}$
- recall that  $V$  is the matrix of eigenvectors (loadings) of  $\hat{\Sigma}$
- order the columns of  $V$  by their corresponding eigenvalues in descending order
- scores:

$$T_{n \times k} = XV$$

- the first column of  $T$  contains the scores for PC1, the second column contains the scores for PC2, etc.

## What is the score?

- PC scores of (standardized)  $X$  are obtained by multiplying  $X$  by the loadings (eigenvectors) of the covariance of  $X$ , say  $\hat{\Sigma}$
- recall that  $V$  is the matrix of eigenvectors (loadings) of  $\hat{\Sigma}$
- order the columns of  $V$  by their corresponding eigenvalues in descending order
- scores:

$$T_{n \times k} = XV$$

- the first column of  $T$  contains the scores for PC1, the second column contains the scores for PC2, etc.

## What is the score?

- PC scores of (standardized)  $X$  are obtained by multiplying  $X$  by the loadings (eigenvectors) of the covariance of  $X$ , say  $\hat{\Sigma}$
- recall that  $V$  is the matrix of eigenvectors (loadings) of  $\hat{\Sigma}$
- order the columns of  $V$  by their corresponding eigenvalues in descending order
- scores:

$$T_{n \times k} = XV$$

- the first column of  $T$  contains the scores for PC1, the second column contains the scores for PC2, etc.

## What is the score?

- PC scores of (standardized)  $X$  are obtained by multiplying  $X$  by the loadings (eigenvectors) of the covariance of  $X$ , say  $\hat{\Sigma}$
- recall that  $V$  is the matrix of eigenvectors (loadings) of  $\hat{\Sigma}$
- order the columns of  $V$  by their corresponding eigenvalues in descending order
- scores:

$$T_{n \times k} = XV$$

- the first column of  $T$  contains the scores for PC1, the second column contains the scores for PC2, etc.



## PCA: example

- data: food consumption in European countries
- figure on the next slide displays the score plot of the first two principal components (scores  $t_1$  and  $t_2$ )
- the score plot is a map of 16 countries: those close to each other have similar food consumption profiles


## PCA: example

- data: food consumption in European countries
- figure on the next slide displays the score plot of the first two principal components (scores  $t_1$  and  $t_2$ )
- the score plot is a map of 16 countries: those close to each other have similar food consumption profiles

## PCA: example


- data: food consumption in European countries
- figure on the next slide displays the score plot of the first two principal components (scores  $t_1$  and  $t_2$ )
- the score plot is a map of 16 countries: those close to each other have similar food consumption profiles

# PCA: example




- score plot illustrates relations between food consumption profiles
- PC1 (PC2) explains 32% (19%) of the variation of the data

# PCA: example




- score plot illustrates relations between food consumption profiles
- PC1 (PC2) explains 32% (19%) of the variation of the data

# PCA: example



- Nordic countries (Finland, Norway, Denmark, and Sweden) are in the upper right corner
- Belgium and Germany are close to the center (origin) of the plot

# PCA: example



- Nordic countries (Finland, Norway, Denmark, and Sweden) are in the upper right corner
- Belgium and Germany are close to the center (origin) of the plot

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data**: subtract the mean of each variable from the corresponding values
- calculate the covariance matrix** of the centered data
- find the eigenvalues and eigenvectors** of the covariance matrix
- choose the principal component(s)**: select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores**: project the centered data onto PCs



## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
  - demean the data**: subtract the mean of each variable from the corresponding values
  - calculate the covariance matrix** of the centered data
  - find the eigenvalues and eigenvectors** of the covariance matrix
  - choose the principal component(s)**: select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
  - calculate the scores**: project the centered data onto PCs

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data:** subtract the mean of each variable from the corresponding values
- calculate the covariance matrix of the centered data
- find the eigenvalues and eigenvectors of the covariance matrix
- choose the principal component(s): select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores: project the centered data onto PCs

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data**: subtract the mean of each variable from the corresponding values
- calculate the covariance matrix** of the centered data
- find the eigenvalues and eigenvectors of the covariance matrix
- choose the principal component(s): select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores: project the centered data onto PCs

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data**: subtract the mean of each variable from the corresponding values
- calculate the covariance matrix** of the centered data
- find the eigenvalues and eigenvectors** of the covariance matrix
- choose the principal component(s)**: select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores**: project the centered data onto PCs

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data**: subtract the mean of each variable from the corresponding values
- calculate the covariance matrix** of the centered data
- find the eigenvalues and eigenvectors** of the covariance matrix
- choose the principal component(s)**: select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores**: project the centered data onto PCs

## PCA: exercise

- dataset:

$x_1$	$x_2$
1	2
3	4
5	6

- perform PCA on this dataset by following these steps:
- demean the data**: subtract the mean of each variable from the corresponding values
- calculate the covariance matrix** of the centered data
- find the eigenvalues and eigenvectors** of the covariance matrix
- choose the principal component(s)**: select the eigenvector(s) associated with the largest eigenvalue(s) as PCs
- calculate the scores**: project the centered data onto PCs

## Exercise: PCA

- demeaned data:

$x_1$	$x_2$
-2	-2
0	0
2	2

- covariance matrix:

$$\begin{pmatrix} 4 & 4 \\ 4 & 4 \end{pmatrix}$$

- eigenvalues and eigenvectors:

- eigenvalues: 0, 8
- eigenvectors:

$$\begin{pmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

$$\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

## Exercise: PCA

- demeaned data:

$x_1$	$x_2$
-2	-2
0	0
2	2

- covariance matrix:

$$\begin{pmatrix} 4 & 4 \\ 4 & 4 \end{pmatrix}$$

- eigenvalues and eigenvectors:

- eigenvalues: 0, 8
- eigenvectors:

$$\begin{pmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

$$\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$



## Exercise: PCA

- demeaned data:

$x_1$	$x_2$
-2	-2
0	0
2	2

- covariance matrix:

$$\begin{pmatrix} 4 & 4 \\ 4 & 4 \end{pmatrix}$$

- eigenvalues and eigenvectors:

- eigenvalues: 0, 8
- eigenvectors:

$$\begin{pmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

$$\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

## Exercise: PCA

- demeaned data:

$x_1$	$x_2$
-2	-2
0	0
2	2

- covariance matrix:

$$\begin{pmatrix} 4 & 4 \\ 4 & 4 \end{pmatrix}$$

- eigenvalues and eigenvectors:

- eigenvalues: 0, 8
- eigenvectors:

$$\begin{pmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

$$\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

## Exercise: PCA

- demeaned data:

$x_1$	$x_2$
-2	-2
0	0
2	2

- covariance matrix:

$$\begin{pmatrix} 4 & 4 \\ 4 & 4 \end{pmatrix}$$

- eigenvalues and eigenvectors:

- eigenvalues: 0, 8
- eigenvectors:

$$\begin{pmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

$$\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$$

## PCA: exercise

- choose the principal component(s): eigenvector associated with the largest eigenvalue (8) is  $\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$  (PC1)
- calculate the scores (projections of the centered data onto PC1):

$$\text{scores} = \text{centered data} \times \text{PC} = \begin{pmatrix} -2\sqrt{2} & 2\sqrt{2} \\ 0 & 0 \\ 2\sqrt{2} & -2\sqrt{2} \end{pmatrix}$$

## PCA: exercise

- choose the principal component(s): eigenvector associated with the largest eigenvalue (8) is  $\begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix}$  (PC1)
- calculate the scores (projections of the centered data onto PC1):

$$\text{scores} = \text{centered data} \times \text{PC} = \begin{pmatrix} -2\sqrt{2} & 2\sqrt{2} \\ 0 & 0 \\ 2\sqrt{2} & -2\sqrt{2} \end{pmatrix}$$

# Shortcomings of standard PCA

- **sensitive to outliers and noise**, which can significantly affect the computed PCs
- **robust PCA** is designed to separate the low-rank structure of the data (true signal) from the sparse noise or outliers
  - more suitable for datasets corrupted by noise or anomalies

# Shortcomings of standard PCA

- **sensitive to outliers and noise**, which can significantly affect the computed PCs
- **robust PCA** is designed to separate the low-rank structure of the data (true signal) from the sparse noise or outliers
  - more suitable for datasets corrupted by noise or anomalies

# Shortcomings of standard PCA

- **sensitive to outliers and noise**, which can significantly affect the computed PCs
- **robust PCA** is designed to separate the low-rank structure of the data (true signal) from the sparse noise or outliers
  - more suitable for datasets corrupted by noise or anomalies



# Robust PCA

- objective: decompose  $X$  into a **low-rank** matrix  $L$  and a **sparse** matrix  $S$  such that  $X = L + S$
- achieved by solving

$$\min_{L,S} \quad ||L||_* + \lambda ||S||_1$$

$$\text{s.t. } X = L + S$$

where  $||L||_*$  is the trace norm of  $L$ ,  $||S||_1$  is the  $L_1$  norm of  $S$ , and  $\lambda$  is a regularization parameter

- $||L||_*$  is the sum of **singular values** (eigenvalues of  $L'L$ )
- $||S||_1$  is the sum of **absolute values** of entries of  $S$

# Robust PCA

- objective: decompose  $X$  into a **low-rank** matrix  $L$  and a **sparse** matrix  $S$  such that  $X = L + S$
- achieved by solving

$$\min_{L,S} \|L\|_* + \lambda \|S\|_1$$

$$\text{s.t. } X = L + S$$

where  $\|L\|_*$  is the trace norm of  $L$ ,  $\|S\|_1$  is the  $L_1$  norm of  $S$ , and  $\lambda$  is a regularization parameter

- $\|L\|_*$  is the sum of **singular values** (eigenvalues of  $L'L$ )
- $\|S\|_1$  is the sum of **absolute values** of entries of  $S$

# Robust PCA

- objective: decompose  $X$  into a **low-rank** matrix  $L$  and a **sparse** matrix  $S$  such that  $X = L + S$
- achieved by solving

$$\min_{L,S} \|L\|_* + \lambda \|S\|_1$$

$$\text{s.t. } X = L + S$$

where  $\|L\|_*$  is the trace norm of  $L$ ,  $\|S\|_1$  is the  $L_1$  norm of  $S$ , and  $\lambda$  is a regularization parameter

- $\|L\|_*$  is the sum of **singular values** (eigenvalues of  $L'L$ )
- $\|S\|_1$  is the sum of **absolute values** of entries of  $S$

# Robust PCA

- objective: decompose  $X$  into a **low-rank** matrix  $L$  and a **sparse** matrix  $S$  such that  $X = L + S$
- achieved by solving

$$\min_{L,S} \|L\|_* + \lambda \|S\|_1$$

$$\text{s.t. } X = L + S$$

where  $\|L\|_*$  is the trace norm of  $L$ ,  $\|S\|_1$  is the  $L_1$  norm of  $S$ , and  $\lambda$  is a regularization parameter

- $\|L\|_*$  is the sum of **singular values** (eigenvalues of  $L'L$ )
- $\|S\|_1$  is the sum of **absolute values** of entries of  $S$

# Clustering

- a bank wants to give credit card offers to its customers; currently, they use customer data to decide which offer should be given to which customer
- the bank can potentially have millions of customers; should it use customer-level data?
- what can the bank do?

# Clustering

- a bank wants to give credit card offers to its customers; currently, they use customer data to decide which offer should be given to which customer
- the bank can potentially have millions of customers; should it use customer-level data?
- what can the bank do?

# Clustering

- a bank wants to give credit card offers to its customers; currently, they use customer data to decide which offer should be given to which customer
- the bank can potentially have millions of customers; should it use customer-level data?
- what can the bank do?

# Clustering

- segment the customers into different groups, e.g. income groups:



- now only three strategies are required, one for each income group
- "high", "average", "low" are not prespecified labels, but outcomes of clustering (**unsupervised learning**)



# Clustering

- segment the customers into different groups, e.g. income groups:



- now only three strategies are required, one for each income group
- "high", "average", "low" are not prespecified labels, but outcomes of clustering (**unsupervised learning**)

# Clustering

- segment the customers into different groups, e.g. income groups:



- now only three strategies are required, one for each income group
- "high", "average", "low" are not prespecified labels, but outcomes of clustering (**unsupervised learning**)

# Clustering

- segment the customers into different groups, e.g. income groups:



- now only three strategies are required, one for each income group
- “high”, “average”, “low” are not prespecified labels, but outcomes of clustering (**unsupervised learning**)

# Clustering

- segment the customers into different groups, e.g. income groups:



- now only three strategies are required, one for each income group
- “high”, “average”, “low” are not prespecified labels, but outcomes of clustering (**unsupervised learning**)

# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation

# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation

# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation

# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation



# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation


# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation


# K-means clustering

- one of the most popular clustering algorithms
- stores  $K$  centroids used to define clusters
- a point is in a cluster if it is closer to that cluster's centroid than any other centroid
- finds the best centroids by alternating between
  - 1 assigning data points to clusters based on current centroids
  - 2 choosing centroids based on the current assignment of data points to clusters
- $K$  is either prespecified or tuned, e.g., by cross-validation


# K-Means




(a)




(b)




(c)



(d)




(e)




(f)

- training examples are dots, cluster centroids are crosses


# K-Means




(a)




(b)




(c)



(d)




(e)



(f)

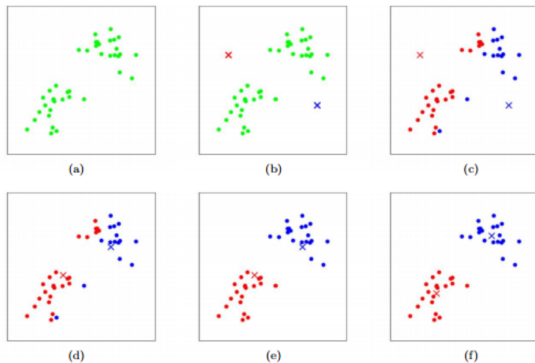
(a) original data

# K-Means




(b) random initialization of cluster centroids

# K-Means



(c)-(f) two iterations of  $K$ -means clustering

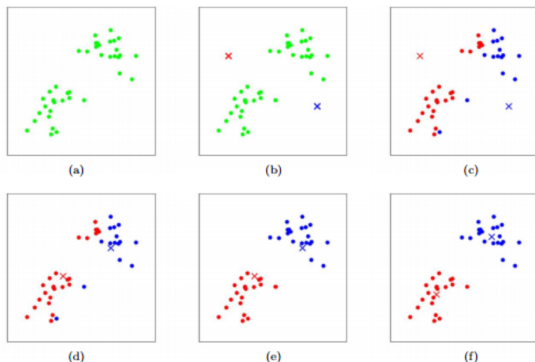
# K-Means



- in each iteration, training examples are assigned to the closest cluster centroid (shown by coloring the training examples with the same color as the cluster centroid to which is assigned)
- then each cluster centroid is moved to the mean of the points assigned to it



# K-Means



- in each iteration, training examples are assigned to the closest cluster centroid (shown by coloring the training examples with the same color as the cluster centroid to which is assigned)
- then each cluster centroid is moved to the mean of the points assigned to it

# K-Means algorithm

- initialize cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
- until convergence, repeat:
  - for each  $i$ , set **membership**

$$c^{(i)} = \operatorname{argmin}_j \|x^{(i)} - \mu_j\|^2$$

- for each  $j$ , set **centroids**

$$\mu_j = \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$

# K-Means algorithm

- initialize cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
- until convergence, repeat:
  - for each  $i$ , set **membership**

$$c^{(i)} = \operatorname{argmin}_j \|x^{(i)} - \mu_j\|^2$$

- for each  $j$ , set **centroids**

$$\mu_j = \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$

# K-Means algorithm

- initialize cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
- until convergence, repeat:
  - for each  $i$ , set **membership**

$$c^{(i)} = \operatorname{argmin}_j \|x^{(i)} - \mu_j\|^2$$

- for each  $j$ , set **centroids**

$$\mu_j = \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$

# K-Means algorithm

- initialize cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^d$  randomly
- until convergence, repeat:
  - for each  $i$ , set **membership**

$$c^{(i)} = \operatorname{argmin}_j \|x^{(i)} - \mu_j\|^2$$

- for each  $j$ , set **centroids**

$$\mu_j = \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}$$

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset

- data:

$$X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$$

- steps:

1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
2. assign points to clusters: assign each point to the nearest centroid
3. update centroids: recalculate the centroids as the mean of the points in each cluster
4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset

- data:

$$X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$$

- steps:

1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
2. assign points to clusters: assign each point to the nearest centroid
3. update centroids: recalculate the centroids as the mean of the points in each cluster
4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset

- data:

$$X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$$

- steps:

1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
2. assign points to clusters: assign each point to the nearest centroid
3. update centroids: recalculate the centroids as the mean of the points in each cluster
4. repeat steps 2 and 3 until convergence



## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset
- data:  
 $X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$
- steps:
  1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
  2. assign points to clusters: assign each point to the nearest centroid
  3. update centroids: recalculate the centroids as the mean of the points in each cluster
  4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset
- data:  
 $X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$
- steps:
  1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
  2. assign points to clusters: assign each point to the nearest centroid
  3. update centroids: recalculate the centroids as the mean of the points in each cluster
  4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset
- data:  
 $X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$
- steps:
  1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
  2. assign points to clusters: assign each point to the nearest centroid
  3. update centroids: recalculate the centroids as the mean of the points in each cluster
  4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- objective: apply K-Means clustering to a small dataset
- data:  
 $X = [(1, 1), (1, 2), (2, 1), (2, 2), (4, 4), (4, 5), (5, 4), (5, 5)]$
- steps:
  1. initialize centroids: choose  $k = 2$  and initial centroids as  $(1, 1)$  and  $(1, 2)$
  2. assign points to clusters: assign each point to the nearest centroid
  3. update centroids: recalculate the centroids as the mean of the points in each cluster
  4. repeat steps 2 and 3 until convergence

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2



## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2



## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

## Exercise: K-Means on a small dataset

- initial centroids:
  - centroid 1:  $(1, 1)$
  - centroid 2:  $(1, 2)$
- iteration 1:
  - cluster 1:  $[(1, 1), (1, 2), (2, 1), (2, 2)]$
  - cluster 2:  $[(4, 4), (4, 5), (5, 4), (5, 5)]$
  - new centroids: centroid 1:  $(1.5, 1.5)$ , centroid 2:  $(4.5, 4.5)$
- iteration 2:
  - no change in cluster assignment
  - convergence achieved
- final centroids:
  - centroid 1:  $(1.5, 1.5)$
  - centroid 2:  $(4.5, 4.5)$
- number of iterations: 2

# Fuzzy C-Means clustering

- $K$  means: each observation only belongs to one cluster
- fuzzy  $C$ -means: each observation may belong to two or more clusters
  - degree of membership of  $x_i$  in cluster  $c$
- often used in pattern recognition and is an extension of the traditional  $K$ -means clustering

# Fuzzy C-Means clustering

- $K$  means: each observation only belongs to one cluster
- fuzzy  $C$ -means: each observation may belong to two or more clusters
  - degree of membership of  $x_i$  in cluster  $c$
- often used in pattern recognition and is an extension of the traditional  $K$ -means clustering

# Fuzzy C-Means clustering

- $K$  means: each observation only belongs to one cluster
- fuzzy  $C$ -means: each observation may belong to two or more clusters
  - degree of **membership** of  $x_i$  in cluster  $c$
- often used in pattern recognition and is an extension of the traditional  $K$ -means clustering

# Fuzzy C-Means clustering

- $K$  means: each observation only belongs to one cluster
- fuzzy  $C$ -means: each observation may belong to two or more clusters
  - degree of **membership** of  $x_i$  in cluster  $c$
- often used in pattern recognition and is an extension of the traditional  $K$ -means clustering

# Fuzzy C-Means clustering

- given a dataset  $X = \{x_1, x_2, \dots, x_n\}$ , FCM partitions the data into  $C$  fuzzy clusters
- each data point  $x_i$  has a degree of membership  $u_{ic} \in [0, 1]$  in each cluster  $c$

# Fuzzy C-Means clustering

- given a dataset  $X = \{x_1, x_2, \dots, x_n\}$ , FCM partitions the data into  $C$  fuzzy clusters
- each data point  $x_i$  has a degree of **membership**  $u_{ic} \in [0, 1]$  in each cluster  $c$



# Fuzzy C-Means clustering

FCM objective function:

$$J(U, V) = \sum_{i=1}^n \sum_{c=1}^C u_{ic}^m \|x_i - v_c\|^2,$$

where

- $U = [u_{ic}]$  is the membership matrix
- $V = \{v_1, v_2, \dots, v_C\}$  are cluster centers
- $m$  is a parameter that controls the level of cluster fuzziness
- $\|x_i - v_c\|$  is the Euclidean distance between the data point  $x_i$  and the cluster center  $v_c$

# Fuzzy C-Means clustering

FCM objective function:

$$J(U, V) = \sum_{i=1}^n \sum_{c=1}^C u_{ic}^m \|x_i - v_c\|^2,$$

where

- $U = [u_{ic}]$  is the membership matrix
- $V = \{v_1, v_2, \dots, v_C\}$  are cluster centers
- $m$  is a parameter that controls the level of cluster fuzziness
- $\|x_i - v_c\|$  is the Euclidean distance between the data point  $x_i$  and the cluster center  $v_c$

# Fuzzy C-Means clustering

FCM objective function:

$$J(U, V) = \sum_{i=1}^n \sum_{c=1}^C u_{ic}^m \|x_i - v_c\|^2,$$

where

- $U = [u_{ic}]$  is the membership matrix
- $V = \{v_1, v_2, \dots, v_C\}$  are cluster centers
- $m$  is a parameter that controls the level of cluster fuzziness
- $\|x_i - v_c\|$  is the Euclidean distance between the data point  $x_i$  and the cluster center  $v_c$

# Fuzzy C-Means clustering

FCM objective function:

$$J(U, V) = \sum_{i=1}^n \sum_{c=1}^C u_{ic}^m \|x_i - v_c\|^2,$$

where

- $U = [u_{ic}]$  is the membership matrix
- $V = \{v_1, v_2, \dots, v_C\}$  are cluster centers
- $m$  is a parameter that controls the level of cluster fuzziness
- $\|x_i - v_c\|$  is the Euclidean distance between the data point  $x_i$  and the cluster center  $v_c$

# Fuzzy C-Means clustering

- iteratively updates membership matrix  $U$  and cluster centers  $V$  until convergence
- membership update:

$$u_{ic}^m = \frac{1}{\sum_{k=1}^C \left( \frac{\|x_i - v_c\|}{\|x_i - v_k\|} \right)^{\frac{2}{m-1}}}$$

- cluster center update:

$$v_c = \frac{\sum_{i=1}^n u_{ic}^m x_i}{\sum_{i=1}^n u_{ic}^m}$$

- stops when changes in membership matrix between consecutive iterations are below a specified threshold

# Fuzzy C-Means clustering

- iteratively updates membership matrix  $U$  and cluster centers  $V$  until convergence
- membership** update:

$$u_{ic}^m = \frac{1}{\sum_{k=1}^C \left( \frac{\|x_i - v_c\|}{\|x_i - v_k\|} \right)^{\frac{2}{m-1}}}$$

- cluster center update:

$$v_c = \frac{\sum_{i=1}^n u_{ic}^m x_i}{\sum_{i=1}^n u_{ic}^m}$$

- stops when changes in membership matrix between consecutive iterations are below a specified threshold

# Fuzzy C-Means clustering

- iteratively updates membership matrix  $U$  and cluster centers  $V$  until convergence
- membership** update:

$$u_{ic}^m = \frac{1}{\sum_{k=1}^C \left( \frac{\|x_i - v_c\|}{\|x_i - v_k\|} \right)^{\frac{2}{m-1}}}$$

- cluster center** update:

$$v_c = \frac{\sum_{i=1}^n u_{ic}^m x_i}{\sum_{i=1}^n u_{ic}^m}$$

- stops when changes in membership matrix between consecutive iterations are below a specified threshold

# Fuzzy C-Means clustering

- iteratively updates membership matrix  $U$  and cluster centers  $V$  until convergence
- membership** update:

$$u_{ic}^m = \frac{1}{\sum_{k=1}^C \left( \frac{\|x_i - v_c\|}{\|x_i - v_k\|} \right)^{\frac{2}{m-1}}}$$

- cluster center** update:

$$v_c = \frac{\sum_{i=1}^n u_{ic}^m x_i}{\sum_{i=1}^n u_{ic}^m}$$

- stops when changes in membership matrix between consecutive iterations are below a specified threshold



# Fuzzy C-Means: example

- dataset with four points: A, B, C, and D
- cluster these points into  $C = 2$  clusters using Fuzzy C-means

- membership values:

Point	Cluster 1	Cluster 2
A	0.8	0.2
B	0.3	0.7
C	0.6	0.4
D	0.1	0.9

# Fuzzy C-Means: example

- dataset with four points: A, B, C, and D
- cluster these points into  $C = 2$  clusters using Fuzzy C-means

- membership values:

Point	Cluster 1	Cluster 2
A	0.8	0.2
B	0.3	0.7
C	0.6	0.4
D	0.1	0.9

# Fuzzy C-Means: example

- dataset with four points: A, B, C, and D
- cluster these points into  $C = 2$  clusters using Fuzzy C-means

- membership values:

Point	Cluster 1	Cluster 2
A	0.8	0.2
B	0.3	0.7
C	0.6	0.4
D	0.1	0.9

## Fuzzy C-Means: example

- **point A** has a high membership value (0.8) in cluster 1 and a low membership value (0.2) in cluster 2, i.e. point A is strongly associated with Cluster 1 but has a slight association with Cluster 2
- **point B** has a membership value of 0.3 in cluster 1 and 0.7 in cluster 2, indicating that it is more closely associated with Cluster 2
- **point C** has a membership value of 0.6 in cluster 1 and 0.4 in cluster 2, suggesting that it belongs more to Cluster 1 but still has some association with Cluster 2
- **point D** has a very low membership value (0.1) in cluster 1 and a high value (0.9) in cluster 2, showing that it is strongly associated with Cluster 2

## Fuzzy C-Means: example

- **point A** has a high membership value (0.8) in cluster 1 and a low membership value (0.2) in cluster 2, i.e. point A is strongly associated with Cluster 1 but has a slight association with Cluster 2
- **point B** has a membership value of 0.3 in cluster 1 and 0.7 in cluster 2, indicating that it is more closely associated with Cluster 2
- **point C** has a membership value of 0.6 in cluster 1 and 0.4 in cluster 2, suggesting that it belongs more to Cluster 1 but still has some association with Cluster 2
- **point D** has a very low membership value (0.1) in cluster 1 and a high value (0.9) in cluster 2, showing that it is strongly associated with Cluster 2

## Fuzzy C-Means: example

- **point A** has a high membership value (0.8) in cluster 1 and a low membership value (0.2) in cluster 2, i.e. point A is strongly associated with Cluster 1 but has a slight association with Cluster 2
- **point B** has a membership value of 0.3 in cluster 1 and 0.7 in cluster 2, indicating that it is more closely associated with Cluster 2
- **point C** has a membership value of 0.6 in cluster 1 and 0.4 in cluster 2, suggesting that it belongs more to Cluster 1 but still has some association with Cluster 2
- **point D** has a very low membership value (0.1) in cluster 1 and a high value (0.9) in cluster 2, showing that it is strongly associated with Cluster 2

## Fuzzy C-Means: example

- **point A** has a high membership value (0.8) in cluster 1 and a low membership value (0.2) in cluster 2, i.e. point A is strongly associated with Cluster 1 but has a slight association with Cluster 2
- **point B** has a membership value of 0.3 in cluster 1 and 0.7 in cluster 2, indicating that it is more closely associated with Cluster 2
- **point C** has a membership value of 0.6 in cluster 1 and 0.4 in cluster 2, suggesting that it belongs more to Cluster 1 but still has some association with Cluster 2
- **point D** has a very low membership value (0.1) in cluster 1 and a high value (0.9) in cluster 2, showing that it is strongly associated with Cluster 2

# Fuzzy C-Means clustering

## Advantages:

- good for overlapping data
- fuzzy membership

## Disadvantages:

- number of clusters should pre-specified
- Euclidean distance may unequally weight underlying factors



# Fuzzy C-Means clustering

## Advantages:

- good for overlapping data
- fuzzy membership

## Disadvantages:

- number of clusters should pre-specified
- Euclidean distance may unequally weight underlying factors